

Reproducibility & project workflow in Python

ZAS Pythonistas working group

May 2026

Table of contents

Session overview	2
Schedule	2
0.1 What is reproducibility?	2
0.2 Virtual environments and package management	3
Virtual environments	3
Basic conda workflow	4
IDEs: VS Code and Positron	4
Exercises: Part A	5
0.3 Project-relative file paths	6
Option 1: <code>pyprojroot</code>	6
Option 2: <code>pathlib</code>	6
Exercises: Part B	7
0.4 Naming conventions & folder structure	8
Recommended folder layout	8
File naming	8
README as a map	8
0.5 Version control with git	9
Core concepts	9
The basic workflow	9
Writing useful commit messages	10
What to put in <code>.gitignore</code>	10
A note on branching	10
Exercises: Part C	10
0.6 Quarto vs. Jupyter notebooks	11
They are not competitors	11
The hidden state problem in notebooks	11
What <code>freeze: true</code> does in Quarto	11
A practical two-stage workflow	11
0.7 R to Python parallels	12
0.8 Resources	12

Session overview

Duration: 90 minutes. **Format:** Discussion + live demo + hands-on exercise.

This session builds from general principles to Python-specific tools – concepts that apply in any language first, then the specific tools and file formats you will use day-to-day in Python.

Schedule

Table 1

[H]

Section	Scope
1. What is reproducibility?	Language-agnostic
2. Virtual environments & file paths	Python-specific
3. Naming conventions & folder structure	Language-agnostic
4. Version control with git	Language-agnostic
5. Quarto vs. Jupyter notebooks	Python-specific
6. Hands-on exercise	Python
7. Wrap-up & next topic vote	–

0.1 What is reproducibility?

Reproducibility means someone else – or future you – can re-run your analysis and get the same result. It operates at two levels:

- **Environment reproducibility:** the same packages and versions are installed. Your code does not break because a collaborator has a different version of pandas.
- **Workflow reproducibility:** the same code runs in the same order and produces the same outputs. No mystery variables, no deleted raw files, no steps that exist only in someone's memory.

Neither layer is sufficient alone. A perfectly locked environment does not help if the script requires you to run cell 7 before cell 3. Clean, ordered code does not help if it requires a package version nobody can install anymore.

Common failure modes: “it works on my machine” (*environment*); raw data was overwritten and the analysis cannot be re-run (*workflow*); a notebook gives different results depending on the order cells were run (*workflow*); a script uses an absolute path like `/Users/name/projects/...` that breaks on every other machine (*workflow*).

Reproducibility is a spectrum – the goal is not perfection from day one, but deliberate choices that reduce friction for your future self and collaborators.

0.2 Virtual environments and package management

Virtual environments

A virtual environment is an isolated Python installation scoped to a single project – the R equivalent of an `renv` library. Do not install packages into your global Python.

This working group uses conda, which is a package manager and virtual environment manager in one. You could also use `venv` (virtual environment manager) + `pip` (package manager) instead for a lighter alternative for pure-Python projects.

Table 2: conda and venv/pip commands are bash code to be entered in the shell (e.g., Terminal, PowerShell, Anaconda Prompt)

[H]

	conda (current setup)	venv + pip (alternative)
Create an environment	<code>conda create -n my-env python=3.11</code>	<code>python -m venv .venv</code>
Install packages	<code>conda install pandas</code> or <code>pip install pandas</code>	<code>pip install pandas</code>
Record env for sharing	<code>conda env export --from-history > environment.yml</code>	<code>pip freeze > requirements.txt</code>
Restore environment	<code>conda env create -f environment.yml</code>	<code>pip install -r requirements.txt</code>
Lockfile (approx. <code>renv.lock</code>)	<code>environment.yml</code> or <code>conda-lock</code> for full pinning	<code>requirements.txt</code> or <code>pyproject.toml</code> + <code>uv.lock</code>
When to use	Non-Python deps, HPC clusters, mixed R+Python	Pure-Python projects, lighter setup

--from-history vs full export: `conda env export` pins every package including transitive dependencies and platform-specific build strings, making it non-portable across operating systems. `--from-history` records only what you explicitly installed, letting conda resolve the rest per platform. For sharing within a team, `--from-history` is usually the better choice.

Avoid conda + pip

If you install packages via `pip` inside conda, `--from-history` won't capture them. Ideally install everything with `conda install` where possible, and use `pip` only for packages not on conda channels.

Why bother with environment-relative package libraries? It's best practice to avoid installing packages into your base environment. Instead, create a separate environment for each project and install only the packages that project needs. This prevents version conflicts between projects — for example, one project requiring an older version of a package than another. It also keeps your setup reproducible: anyone can recreate your exact environment from your `environment.yml`. It's the same logic behind using Rprojects (environment) and `renv` lockfiles (package management) when working in R.

A note on terminals: Mac vs Windows

Most code in this handout is written for **bash** (the shell used on macOS and Linux). Windows uses PowerShell by default, which has different syntax. Windows users should use one of the following so the commands here work as written:

Table 3: Terminal options

[H]

Option	Description	For this session
Git Bash	Installs with Git for Windows. Gives a full bash shell. Simplest option for this session.	Yes – simplest
WSL (Windows Subsystem for Linux)	A full Linux environment inside Windows. More powerful but requires one-time setup.	Optional
Positron / VS Code integrated terminal	Auto-uses bash or Git Bash if configured. Conda commands work out of the box.	Yes
Anaconda Prompt	A Windows terminal pre-configured for conda. Conda commands work but bash syntax may not.	Conda only

💡 Conda in IDE integrated terminal

You can use conda in the Positron / VS Code integrated terminal on macOS without any additional setup. On Windows, first run `conda init powershell` from Anaconda Prompt, then close and reopen your IDE. You only need to do this once.

Basic conda workflow

Listing 1 Terminal (bash / Git Bash / WSL)

```
conda env list                # list all environments
conda create -n zas-python-wg python=3.11  # create once
conda activate zas-python-wg             # activate
conda info --envs                 # check active environment
conda install pandas matplotlib        # install packages
conda env export --from-history > environment.yml # record environment
conda env create -f environment.yml     # restore (others)
```

Positron / VS Code: Command Palette -> *Python: Select Interpreter* -> choose your conda environment. The integrated terminal will use your active conda environment automatically.

IDEs: VS Code and Positron

Both VS Code and Positron are built on the same engine (VS Code's core), so the interface is nearly identical. Positron is Posit's fork tuned for data science: it adds a persistent Variables pane, a Console panel, and a Plots viewer – all familiar from RStudio. For this working group either IDE works; Positron is recommended if you are coming from RStudio.

Table 4

[H]

	VS Code	Positron
Built on	VS Code core	VS Code core (Posit fork)
Python interpreter selection	Command Palette -> Python: Select Interpreter	Same – Command Palette -> Python: Select Interpreter
Environment activation	Auto-activates <code>.venv</code> or <code>conda env</code> in workspace	Same – auto-activates on project open
Variables pane	Via Python extension (limited)	Built-in, always visible (like RStudio)
Console panel	Integrated terminal only	Built-in, persistent (like RStudio)
Plots viewer	Via Jupyter extension	Built-in viewer pane
Notebook support	Full <code>.ipynb</code> support	Full <code>.ipynb</code> support
R support	Via R extension	First-class, built-in
Recommended for	General development, large teams	Data science, R users moving to Python

! Your typical workflow

1. Open Positron (or VS Code)
2. **File -> Open Folder** -> navigate to your project folder
3. Check the interpreter in your **status bar** – it should show your conda environment name (e.g. `zas_python_wg`). If not: Command Palette -> Python: Select Interpreter -> choose it. Alternatively, check in Terminal with `which python` and `python --version` (but you can't select it in the Terminal).
4. Open the integrated terminal ('Ctrl+' or Terminal -> New Terminal) – it should activate your conda environment automatically
5. Verify with `conda info --envs` (asterisk = active env) and `pwd` (correct folder). If there is no asterisk (*) next to your environment, run `conda activate zas_python_wg` and

verify.

Exercises: Part A

Open the `repro_exercises.py` file and do the exercises under Part A 'Check your environment'.

0.3 Project-relative file paths

Hardcoded absolute paths break on every machine except the one they were written on. There are two good ways to fix this in Python.

Option 1: `pyprojroot`

This is the recommended option as it's closest to `here : : here ()` in R.

Install once: `conda install pyprojroot` (or `pip install pyprojroot` if not using conda)

Listing 2 Python

```
from pyprojroot import here
from pathlib import Path

data_path = here("data/raw/stimuli.csv")
output_path = here("outputs/figures")
```

`pyprojroot` walks up from the current file until it finds a project root marker – `.git`, `pyproject.toml`, `environment.yml`, etc. – and anchors all paths there. This is the direct Python equivalent of `here : : here ()` and works correctly regardless of where the script lives in the project.

Updating `environment.yml`

Whenever you install a new package, update your `environment.yml` so others can reproduce your environment:

Listing 3 Terminal

```
conda env export --from-history > environment.yml
```

Whenever you open a project, remember to activate its environment before running any code:

Listing 4 Terminal

```
conda activate <env-name>
```

This ensures you're using the right packages (even on your own machine) which is also why keeping `environment.yml` up to date matters, since that's what recreates the environment elsewhere.

Option 2: `pathlib`

This option requires no extra installation.

This works without any extra packages but requires counting `.parent` calls correctly – which breaks if you move the script to a different subfolder depth. Use `pyprojroot` if you want `here ()`-style convenience.

Listing 5 Python

```
from pathlib import Path

# __file__ is the path of the current script
# .parent navigates up one folder at a time -- count carefully
PROJECT_ROOT = Path(__file__).parent.parent

data_path = PROJECT_ROOT / "data" / "raw" / "stimuli.csv"
output_path = PROJECT_ROOT / "outputs" / "figures"
output_path.mkdir(parents=True, exist_ok=True) # create if needed
```

Exercises: Part B

Open the `repro_exercises.ipynb` file and do the exercises under Part B 'Fix the broken file path'. Try the same exercises in `repor_exercises.qmd`. How do you find working in one versus the other?

0.4 Naming conventions & folder structure

These conventions apply in any language. Consistent structure makes projects navigable by collaborators and comprehensible to your future self.

Recommended folder layout

Listing 6 Project structure

```
my-project/  
  data/  
    raw/           # original data -- never modified  
    processed/     # derived data, reproducible from raw  
  scripts/         # numbered analysis scripts  
  notebooks/       # exploratory .ipynb files  
  outputs/  
    figures/  
    tables/  
  docs/            # manuscripts, protocols, notes  
  environment.yml  # conda environment (or requirements.txt)  
  .gitignore  
  README.md
```

! Raw data is sacred

Never overwrite raw data and never modify it in place. All transformations happen in code and produce outputs in data/processed/.

File naming

Table 5

[H]		
Rule	Avoid	Prefer
Lowercase only	My Analysis.py	my-analysis.py
No spaces	model results final.csv	model-results-final.csv
Date-prefix versioned out-puts	results.csv	2025-05-04_model-results.csv
Verb-first for scripts	data.py	01_clean-data.py
Zero-pad numbered scripts	1_clean.py, 10_model.py	01_clean.py, 10_model.py
Be specific	analysis.py, data2.csv	01_clean-eyetracking.py

README as a map

Every project should answer three questions: what does this project do, how do I run it, and what do the outputs mean?

Listing 7 README.md

```
# Project name
One sentence describing the study or task.

## Setup
conda env create -f environment.yml
conda activate zas-pythonistas

## Running the analysis
python scripts/01_clean-data.py
python scripts/02_fit-model.py

## Outputs
outputs/figures/  -- all plots, referenced in the manuscript
outputs/tables/   -- model results tables
```

0.5 Version control with git

Git tracks changes to your files over time, so you can see what changed, when, and why – and roll back if something breaks. It is language-agnostic and works for any text-based project. GitHub (and GitLab, Codeberg, etc.) are hosting platforms for git repositories; git itself runs entirely on your local machine.

Core concepts

- **Repository (repo):** a folder tracked by git, containing your project and its full history
- **Commit:** a saved snapshot of your project at a point in time, with a message describing what changed
- **Staging area:** a holding zone where you choose exactly which changes to include in the next commit – you do not have to commit everything at once
- **Branch:** a parallel version of the project where you can develop a feature or fix without affecting the main version (more on this below)

The basic workflow

Listing 8 Terminal

```
# --- One-time setup ---
git config --global user.name "Your Name"
git config --global user.email "you@example.com"

# --- Start tracking a project ---
git init                                # initialise a new repo in the current fol

# --- Day-to-day workflow ---
git status                             # see what has changed
git add scripts/01_analyse.py          # stage a specific file
git add .                              # stage everything (use carefully)
git commit -m "fix: use relative paths in analysis script"

# --- Review history ---
git log --oneline                      # compact list of past commits
git diff                              # see unstaged changes line by line
```

Writing useful commit messages

A good commit message completes the sentence “If applied, this commit will...”. Use a short imperative verb phrase: `add`, `fix`, `update`, `remove`. Keep it under 72 characters.

Table 6

[H]

Avoid	Prefer
<code>fixed stuff</code>	<code>fix: use relative paths in analysis script</code>
<code>update</code>	<code>add: conda environment file and README setup instructions</code>
<code>changes</code>	<code>remove: hardcoded absolute paths from all scripts</code>
<code>asdfgh</code>	<code>refactor: split data cleaning into separate script</code>

What to put in `.gitignore`

A note on branching

Branches let you work on a new feature or experiment without touching the main version of your project. When the work is ready, you merge it back. For solo projects, you may not need branches at all – a clean linear commit history is fine. For collaborative projects, a simple convention is: keep `main` stable and do new work on a named branch.

We will cover branching and collaborative workflows (pull requests, merge conflicts) in a future session.

Exercises: Part C

Open the `repro_exercises` exercises in whichever file format you prefer (`.py`, `.ipynb`, `.qmd`) and do the exercises under Part A ‘git hygiene’.

0.6 Quarto vs. Jupyter notebooks

They are not competitors

Quarto *uses* Jupyter as its execution engine – adding `jupyter: python3` to a `.qmd` header tells Quarto to call Jupyter to run your code. The real choice is between two **file formats**:

Table 7

[H]

	.ipynb notebook	.qmd (Quarto)
File format	.ipynb (JSON)	.qmd (plain text)
Execution order	Any order – cells run independently	Always top-to-bottom, fresh process
Kernel state	Persists between cells; can be stale	No persistent state between renders
Git diffs	Noisy – outputs embedded in JSON	Clean – plain text only
Output storage	Stored inside the file	Separated from source
Multi-format export	Requires nbconvert + manual steps	Built-in: HTML, Word, PDF, slides
Inline computed values	Not straightforward	<code>'python round(val, 2)'</code>
Best for	Exploration, iteration, teaching	Reports, papers, shared outputs

The hidden state problem in notebooks

The biggest reproducibility risk in `.ipynb` files is **out-of-order execution**. Jupyter remembers every variable set during a session, even from cells you have deleted or moved. A notebook that appears to work may depend on a variable set hours ago in a cell that no longer exists. The only reliable check is *Kernel -> Restart & Run All* before sharing – and people often forget.

What `freeze: true` does in Quarto

By default, Quarto re-runs all your code from scratch on every render – correct for reproducibility, but slow for heavy computations like model fitting or EEG preprocessing. `freeze: true` changes this: Quarto **caches each chunk's output** and only re-runs a chunk when its source code actually changes.

- **First render:** runs everything, saves outputs into a `_freeze/` folder alongside your `.qmd`
- **Subsequent renders:** skips unchanged chunks and reuses saved outputs
- **After editing a chunk:** only that chunk re-runs

Committing `_freeze/` to git means collaborators can render the document and see your outputs without needing your data or compute environment.

R parallel: `freeze: true` is roughly analogous to `knitr cache = TRUE` in R Markdown, but more reliable – it caches at the file level, tracks source changes explicitly, and plays well with git.

A practical two-stage workflow

Use notebooks for the messy exploratory work. Once the approach is settled, move clean code into a `.qmd` for the shareable output.

0.7 R to Python parallels

Table 8

[H]

Concept	R	Python
IDE	RStudio / Positron	Positron / VS Code
Isolated environment	renv library (.Rprofile)	conda env or .venv (venv)
Package install	install.packages() / pak::pak()	conda install / pip install
Dependency lockfile	renv.lock	environment.yml or requirements.txt or uv.lock
Restore environment	renv::restore()	conda env create -f environment.yml or pip install -r requirements.txt
Project root	.Rproj file sets working directory	No automatic root – use explicit Path(__file__) convention
Project-relative paths	here::here()	here("data/raw/file.csv") via pyprojroot or Path(__file__).parent / ...
Set random seed	set.seed(42)	random.seed(42) + numpy.random.seed(42)
Dynamic report format	.Rmd or .qmd	.qmd with jupyter: python3
Freeze / cache computations	knitr cache = TRUE	freeze: true in Quarto
Inline computed value	'r round(val, 2)'	'python round(val, 2)'
Version control	git (same)	git (same)

0.8 Resources

- **pythRon book** (ZAS internal): R/Python concepts side-by-side with worked examples
- **Quarto + Python:** quarto.org/docs/computations/python.html
- **Quarto freeze:** quarto.org/docs/projects/code-execution.html#freeze
- **uv:** docs.astral.sh/uv/
- **Pro Git book (free):** git-scm.com/book/en/v2
- **Real Python – virtual environments primer:** realpython.com/python-virtual-environments-a-primer
- **The Turing Way** – community handbook for reproducible research: the-turing-way.netlify.app

Listing 9 .gitignore

```
# Python
.venv/
__pycache__/
*.pyc
*.pyo
.ipynb_checkpoints/

# Conda
.conda/

# Quarto
/.quarto/
_book/
*_files/

# R
.Rhistory
.RData
.Rproj.user/

# LaTeX auxiliary files
*.aux
*.log
*.out
*.toc
*.fls
*.fdb_latexmk
*.synctex.gz

# Temporary conda exports
environment-full.yml
environment-minimal.yml

# Secrets -- never commit these
.env
*.key

# macOS
.DS_Store

# Large data files (store on OSF or NAS)
data/raw/
*.edf
*.mat
```

Listing 10 Terminal

```
git branch feature/add-preprocessing # create a branch
git checkout feature/add-preprocessing # switch to it
                                         # (or: git checkout -b name, to do both)
git checkout main                      # switch back to main
git merge feature/add-preprocessing   # merge when ready
```

Listing 11 _quarto.yml

```
execute:  
  freeze: true
```

Listing 12 Terminal

```
# Extract clean code from a notebook when you are ready  
jupyter nbconvert --to script exploration.ipynb  
# produces exploration.py -- paste the useful parts into your .qmd
```
